

Knowledge and Principles

Guide 2 of 6

TL;DR

Theseus turns uploaded material into firm principles through a four-stage pipeline: ingest, claim extraction, clustering, and principle distillation. The first two stages run automatically; the last two involve a triage queue you sit in front of when you want to keep the firm's working memory clean. This guide walks the pipeline end to end, shows where each artifact lives, and explains the founder's role in triage.

You will need

- A signed-in Codex session.
- At least one upload in the corpus — if you do not have one, complete Guide 1 first.
- Optional: a couple of related documents already uploaded, so that clustering has something to merge across.

This guide does NOT cover

- Question-answering with the Oracle — Guide 3.
- Currents opinions that cite principles — Guide 4.
- Forecasts that may use principle-backed conclusions as a reasoning step — Guide 5.
- Operator-only deletion / retraction of principles — Guide 6.

1 The pipeline in one diagram

Upload → Claims → Clusters → Principles
ingest layer → claim extractor → clustering pass → distillation

Each arrow above corresponds to a worker stage. Stages 1 and 2 run on upload. Stages 3 and 4 run on a schedule and are also triggerable from the founder triage queue.

2 Stage 1 — ingest

Uploads enter through **Upload** on the Codex (/ (authed) /upload). The ingest worker normalizes the file (MIME detection, encoding repair, page extraction for PDFs), persists it to storage, and writes an **Upload** row with `publishedAt=NULL`, `visibility='private'`, and a processing badge that the UI watches.

What you see in the Codex: the new row appears in **Upload** list, then in **Knowledge** once at least one claim has been extracted from it. Provenance is captured at ingest: file hash, source URL (if any), uploader identity, and timestamps.

3 Stage 2 — claim extraction

The claim extractor reads the normalized upload and emits short, declarative, falsifiable statements. Each claim carries:

- the upload it came from (provenance);
- a topic label and an embedding vector for clustering;
- the verbatim span it summarizes (so the Oracle can cite it);
- an extractor-confidence score.

Look at any claim in **Knowledge** and you can click straight back to the source line. This is the contract that makes everything downstream auditable: no claim is allowed to exist without a pointer to the text that produced it.

[screenshot screenshots/02_knowledge_claim_view.png not yet captured — The claim view in Knowledge, with a single claim expanded to show its source span.]

4 Stage 3 — clustering

The clustering pass groups claims by topic and semantic similarity. The output is a candidate cluster: a small set of claims that look like they are talking about the same thing.

Clusters are not yet principles. They are an intermediate object: an invitation to consider whether the firm wants to lift the underlying material into a reusable principle.

5 Stage 4 — principle distillation

For each candidate cluster the distillation worker produces a draft **Principle**: a single, sharp statement that the underlying claims support, plus an explicit list of the supporting claims and an explicit list of the strongest objections found in the corpus.

A principle has three states that matter:

- `draft` — not visible publicly. Visible in the triage queue.
- `accepted` — vetted. May or may not be public.
- `rejected` / `merged` / `needs_rereview` — terminal or paused states.

A principle becomes publicly visible only when `status = 'accepted'` **and** `publicVisible = true`. Both are required. See `docs/operator/public_surfacing.md` for the authoritative predicate table.

6 The founder's role — the triage queue

The principles surface in the Codex (`/ (authed) /principles`) is the triage queue. Open it and you will see candidate principles in three columns:

- **New drafts.** Just produced by distillation. Read the principle, scan supporting claims, scan objections. Decide: `accept`, `merge into existing`, or `reject with reason`.
- **Needs re-review.** Either contradicted by a newer cluster or flagged by the adversarial pass. The previous decision should be re-justified or rolled back.
- **Accepted.** The firm's working memory. Editable in place; toggling `publicVisible` sends a principle to (or removes it from) the public methodology page.

[screenshot screenshots/02_principles_triage.png not yet captured — The triage queue, with the three columns described above.]

The triage queue is the only stage in this pipeline that requires human judgment. If you ignore it for a week, drafts accumulate; the firm's automatic answers keep working off the last batch of accepted principles. Nothing collapses, but the firm's public surface stops evolving.

7 Where principles live afterwards

- **Codex (private).** `/ (authed) /principles`. The full draft + accepted + rejected set, with editing controls.
- **Public methodology page.** `/methodology/principles` on the public site. The accepted-and-public subset.
- **Currents opinions.** An opinion may cite principles when it leans on a firm position rather than a one-off claim.
- **Oracle answers.** The Oracle may cite principles for the same reason.

8 Common failure modes

- **Drafts pile up.** The clustering pass keeps proposing principles; nobody reviews them. Reserve a recurring slot weekly for triage. Five to ten minutes per draft is realistic.
- **A principle on the public site contradicts a private accepted one.** Flip `publicVisible=false` on the public one and re-triage. Both stay in the corpus; only their public surfacing changes.
- **The cluster is junk.** Reject the draft with a reason. The reason is read by the next distillation pass, which uses it to avoid re-proposing the same cluster.

9 Where the docs are

- `docs/methods/MQS_Specification.md` — the formal claim / conclusion / principle specification.
- `docs/operator/public_surfacing.md` — which fields gate each public surface, including principles.
- `noosphere/noosphere/distillation/` — the worker code, for the curious.

Where to read next

Continue with [Guide 3 — The Oracle](#) for how question-answering uses the corpus you have just helped curate.