

Knowledge and Principles

Guide 2 of 6

TL;DR

Everything the firm reads, hears, or records eventually becomes a *claim* (one sentence, one speaker), gets grouped with similar claims into a *principle* (a durable belief), and earns a *conclusion* (a vetted, principle-shaped position) when the workshop is willing to commit. This guide walks the full corpus side of the Codex: the surfaces (**Library, Upload, Transcripts, Knowledge, Conclusions, Contradictions, Open Questions, Explorer, Knowledge graph, Principles, Methods**), what each is for, and the small set of routine actions you take to keep the firm's working memory clean. Approval, deletion, and retraction of public material is operator-only and is covered in Guide 6.

You will need

- A signed-in Codex session.
- At least one ingested upload — if you do not have one, complete Guide 1 first.
- Twenty to thirty minutes the first time through; about ten minutes a week for routine triage afterwards.

This guide does NOT cover

- Question-answering and the Research Advisor — Guide 3.
- Currents opinions that cite conclusions — Guide 4.
- Forecasts that use conclusions as a reasoning step — Guide 5.
- The publication-review checklist, source-standing changes review, deletion vs. retraction, and the methodology review week (all operator-only) — Guide 6.

1 One picture of the corpus side

Upload → Claims → Clusters → Principles
extract → embed → distill → triage

Claims → Conclusions → Publication
methodology profile → adversarial cycle → public surface

Two passes happen on every ingested upload. The top row distills durable *principles* from clustered claims. The bottom row builds *conclusions* — principle-shaped, source-grounded positions that the firm is willing to commit to, accompanied by a methodology profile (how this source reasoned) and adversarial objections (what the strongest case against this conclusion looks like).

2 The Library — everything the firm has uploaded

Library (/library) is the org-wide inventory of every upload. Rows include written documents, podcast transcripts, Dialectic session files, annotation bundles, and external material. Each row shows the title, the source type, the founder who uploaded it, the visibility (org, semi-private, or private), and the processing state.

You can delete your own uploads from this page; the file goes through a short “please confirm” step and is then soft-deleted. If you need something somebody else uploaded removed, the library lets you open a **Deletion request** that the owner accepts or declines. (Outright hard deletion of material the firm has already cited publicly is operator-only and goes through a different surface — see Guide 6.)

3 Upload — raw material in

Upload (/upload) is the form. Drop a file in, pick a source type, pick a visibility, optionally check “publish as a blog post.”

The form accepts text-like files (.md, .txt, .vtt, .jsonl, .json), PDFs, .docx, and audio (.mp3, .m4a, .wav, .webm, .ogg). For very large files, the Codex automatically uses a direct-to-storage upload flow behind the scenes.

After submit, the workshop runs in the background. The detail page for your upload streams the workshop’s progress live — you can refresh and see fresh lines. When the workshop finishes, the badge clears, the upload’s status becomes *ingested*, and counts appear next to the row: how many claims were extracted, how many clusters were produced, how many methodology profiles were written.

[screenshot screenshots/02_upload_processlog.png not yet captured — An upload’s detail page during processing, with the workshop’s progress streamed live.]

4 Transcripts — reading the source

For uploads with a transcript shape (podcasts, Dialectic sessions, VTT files), **Transcripts** (/transcripts/<upload>) presents the text in chunks, with speaker labels and millisecond timestamps where those are available. This is the source-of-truth view: every claim elsewhere in the system carries a pointer that lands here, on the exact span it came from.

5 The atomic claim

A claim is one sentence, one speaker, one assertion, plus a fingerprint of its meaning. The workshop refuses to roll claims into paragraphs — the whole point is to be able to reason about them individually.

What a claim carries:

- The text.
- Who said it (speaker label, founder identity if known).
- Where it came from (upload, source span start and end).
- Its disciplines (one or more, from a fixed vocabulary).
- Its fingerprint (an embedding vector).
- A type — factual, methodological, normative, predictive, definitional, or interpretive.

- Its origin — `founder` (the speaker’s own assertion) or `external` (a quoted external position; the workshop strips these out of inference).
- Hedges and evidence pointers.

The workshop runs a guard on input that strips out interview prompts and quoted external positions *before* segmentation, so “what do you think of X?” (the host asking) doesn’t end up as the founder asserting X.

6 Knowledge — claims, browsable

Knowledge (/knowledge) is the claim browser. Filter by upload, by speaker, by discipline, by claim type. Click any claim to see its source span, its embedding neighbors, and the conclusions that cite it.

7 Conclusions — the firm’s vetted positions

Conclusions (/conclusions) is the most important surface on the corpus side. A conclusion is a principle-shaped, source-grounded statement of something the firm is willing to commit to.

Every conclusion carries:

- Its text and rationale.
- A **confidence tier** on a ladder: `firm` (most confident), `founder` (a founder vouches), `open` (under review), or one of the demoted tiers (`moderate`, `low`, `speculative`) that the adversarial cycle can drop a conclusion into.
- Its supporting principles, its evidence chain (the underlying claims), and any dissenting claims.
- Its source uploads.
- Its **methodology profile** — a separate record of *how* the source reasoned (see below).
- Its **quality score** (MQS) on the five working criteria (progressivity, severity, aim-method fit, compressibility, domain sensitivity).
- Its **principle shape** contract (next section).

8 The principle-shape contract

The firm requires every principle-grade conclusion to declare what kind of thing it is and where it applies. The contract has five mandatory fields:

- `principleKind` — one of seven shapes: `RULE` (a normative “do X”), `CRITERION` (“X is the test for Y”), `MECHANISM` (“X causes Y by Z”), `HEURISTIC` (“X usually works in case Y”), `DEFINITION`, `FORMULA`, or `ALGORITHM`.
- `domainOfApplicability` — a 300-character description of *where this is meant to hold and where it stops working*. Required. “Always” is not acceptable.
- `quantifiableProxies` — up to five measurable proxies that would let an outside reviewer falsify the principle.
- `decisionExamples` — up to three concrete decisions the principle would direct.
- `sourceSpan` — a verbatim substring of a source chunk that the principle is anchored to. If the quoted span doesn’t literally appear in the source, the workshop refuses to write the row.

This is the firm’s standing answer to “where would this stop being true?” A conclusion without a domain of applicability cannot graduate to `firm` tier.

[screenshot screenshots/02_conclusion_detail.png not yet captured — A conclusion’s detail page: text on the left; principle-shape contract, sources, MQS pill, and adversarial status on the right.]

9 Quantitative Formalisation — the falsifiability layer

For any conclusion that is being prepared for publication, the quantitative formalisation page (`/principles/<id>/quantitative`) shows whether a falsifiability layer has been attached: a null hypothesis, one or more metrics, one or more statistical tests, and one or more data sources. The workshop refuses to mark a quantitative formalisation `APPROVED` without all four.

This is the operational form of the working criterion “the test should be severe enough that the claim would have failed if it were wrong.” If you cannot state what would falsify the principle, the principle does not get to be public.

10 Contradictions and Open Questions

Contradictions (`/contradictions`) lists active claim-vs-claim disagreements the workshop has flagged. Each row carries: claim A, claim B, the severity score, a six-layer breakdown of how the coherence engine voted, a short narrative, and a status (`active`, `resolved`, or `dismissed`).

Open Questions (`/open-questions`) is the unresolved bin: pairs the six layers couldn’t reach a majority on. These are good candidates for the next conversation to revisit.

Resolve or dismiss either kind from its detail page, attaching a brief note. The note is part of the audit trail.

11 Explorer — the embedding map

Explorer (`/explorer`) is a 3D map of the firm’s conclusions in embedding space. Each conclusion is a point; similar conclusions cluster together. The view supports:

- Saved camera positions encoded in the URL, so you can share a view by copy-pasting the link.
- Hotkeys for toggling labels, axes, and confidence tiers.
- Empty-state diagnostics — if the embedding index isn’t built for some conclusions, the panel tells you which ones.

[screenshot screenshots/02_explorer.png not yet captured — The Explorer view, with conclusions as points and confidence tiers color-coded.]

A second, graph-shaped view lives at **Knowledge graph** (`/knowledge`, the graph tab) and shows principles as nodes with typed edges (`SUPPORTS`, `CONTRADICTS`, `REFINES`, `INSTANTIATES`, and others).

12 The principle triage queue

Principles (`/principles/queue`) is the queue of candidate principles the workshop has distilled but not yet committed to. Three columns:

- **New drafts.** Just produced. Read the principle text, scan supporting claims, scan objections. Decide: `accept`, `merge into existing`, or `reject with reason`.

- **Needs re-review.** Either contradicted by a newer cluster or flagged by the adversarial pass. The previous decision should be re-justified or rolled back.
- **Accepted.** The firm's working memory. Editable in place. Whether an accepted principle is also *publicly visible* is a separate operator-only step.

Principles → **Triage** (`/principles/<id>/triage`) is where you commit a decision. The reason you give for a rejection is read by the next distillation pass, which uses it to avoid re-proposing the same cluster.

13 Methodology profiles — how a source reasons

Methods (`/methods`) is the firm's registry of reasoning patterns: `first_principles_decomposition`, `adversarial_revision`, `analogical_transfer`, `dialogic_unfolding`, and others. Each method has a one-sentence rule, a transfer condition (when to use it), an assumption it presumes, and a known failure mode.

For every upload, the workshop writes a *methodology profile*: a short record of *how* this piece of material reasoned, with four verbatim anchor quotes pinned to sentence indices in the source. The profile carries a confidence score, the inferred transfer targets, and a dedupe key so re-running over the same material updates rather than duplicates.

Methods → `[name]/[version]` surfaces a per-method detail page that includes:

- A track record: per-domain Brier score, calibration slope with bootstrap CI, and the severity-pass rate.
- A failures table.
- A domain bound (the centroids that define the method's domain of applicability in embedding space).
- A changelog (every version of the method source).

A method can be `active`, `under_review`, `deprecated`, or `retired`. **Methods** → **Candidates** (`/methods/candidates`) is the inbox of proposed new methods; the **Methods graph** (`/methods/graph`) shows their dependency structure.

14 Source standing — when a citation's status changes

The firm's published material cites outside sources. Those sources can change status over time — a paper gets retracted, a study is corrected, an article is disputed, an article expires behind a paywall. The workshop tracks these as *source-standing* events. When a standing flips, every conclusion that cited the affected source is flagged for review.

You will see the resulting attention items on your dashboard. The review action itself — confirming, overriding, or dismissing the flag — is operator-only and is covered in Guide 6. For most founders, the impact of source-standing changes shows up as conclusions that quietly get an updated status banner.

15 Cascade — what would change if I revised this?

Cascade (`/cascade/<conclusionId>`) is the impact-analysis view. Pick a conclusion; the page shows every downstream artifact whose production graph touches it: which methods produced what, which methodology profiles depend on it, which Currents opinions and Forecasts cite it. This

is what makes the firm’s “reanalyze without wiping” rule operational: when a method version bumps, the cascade is the precise set of things to revisit.

16 Decay — staleness

Decay (/decay) lists the firm’s knowledge that is aging or stale. A conclusion ages by a fixed schedule or by trigger:

- **Fixed interval** — e.g., revisit annually.
- **Evidence changed** — a source the conclusion depends on was retracted or its standing changed.
- **Method version bump** — the method that produced the conclusion has a new version.
- **Embedding drift** — the conclusion’s meaning has shifted beyond a threshold relative to its origin snapshot.
- **Outcome observed** — a forecast or prediction the conclusion implied has resolved.
- **Calibration regression** — the method’s per-domain track record has worsened beyond a threshold.

Each row offers a revalidate, retire, or update action. Revalidating re-runs the relevant workshop passes; retiring marks the conclusion `retired` but keeps the row for audit; updating creates a new version.

17 The founder’s weekly rhythm on the corpus side

- **Daily glance. Dashboard.** See active principles, recent conclusions, drift events, operations health.
- **15 minutes weekly. Principles triage queue.** Accept, merge, or reject the drafts that have accumulated.
- **10 minutes weekly. Contradictions and Open Questions.** Resolve or dismiss with a note.
- **As needed. Decay** when something aged; the attention items on the dashboard when they appear.

If you ignore triage for a week, drafts accumulate but nothing collapses — the firm’s automatic answers just keep working off the last batch of accepted principles. The corpus stops evolving, not the system.

18 Common failure modes

- **Drafts pile up.** Reserve a recurring slot for triage. Five to ten minutes per draft is realistic.
- **The cluster is junk.** Reject the draft with a reason. The reason feeds back into the next distillation pass.
- **A conclusion was demoted by the adversarial cycle.** Read the objection. Decide whether the firm has a counter you’d like to surface, or whether the demotion is just.
- **A source you depended on got retracted.** The system already filed an attention item; founder-alpha will work through it on the operator side.

19 Where the docs are

- `docs/methods/MQS_Specification.md` — the formal claim / conclusion / principle specification, including the principle-shape contract.
- The other five guides in this series.

Where to read next

Continue with [Guide 3 — The Oracle](#) for how question-answering uses the corpus you have just helped curate, including the Research Advisor and the Reading Queue.