

The Oracle, Research Advisor, and Reading Queue

Guide 3 of 6

TL;DR

The Oracle is the firm's question-answering surface. It does not run a free chat against a foundation model; it answers *from* the corpus the workshop has built and refuses — or warns — when the corpus is too thin. This guide explains the citation contract (every quoted span must literally appear in the source), the confidence bands, the hallucination warning, and what to do with a low-confidence answer. It also covers two adjacent surfaces: the **Research Advisor** (what to read next, and why) and the **Reading Queue** (your personal stack).

You will need

- A signed-in Codex session.
- A working corpus — at minimum a handful of uploads and a few accepted principles, otherwise the Oracle will mostly abstain.

This guide does NOT cover

- How claims, principles, and conclusions get into the corpus — Guide 2.
- Currents opinions, which use a different (event-grounded) prompt path — Guide 4.

1 What the Oracle is

The Oracle lives at `/ask` (public, the firm's open Q&A) and at `/codex-ask` (founder-facing, with access to private material). Both accept a question in natural language and return:

- an answer, written in the firm's voice;
- inline citations to claims, conclusions, principles, or uploads;
- a confidence band: `high`, `medium`, `low`, or `abstain`;
- a hallucination-risk badge when the answer text contains a span the citation chain does not unambiguously support.

It does not return: free-form essays not grounded in the corpus, opinions on outside events (that is Currents' job), or speculative projections (that is the forecasts portfolio's job).

[screenshot screenshots/03_oracle_answer.png not yet captured — An Oracle answer with inline citations, confidence band, and hallucination badge.]

2 How the Oracle composes an answer

The path under the hood is short and conservative. The Oracle:

1. Encodes your question into a fingerprint.

2. Retrieves the top-ranked principles and conclusions whose fingerprints are nearest to the question's, filtered by a minimum conviction threshold.
3. Pulls the supporting claims for each retrieved principle.
4. Hands the principles (as axioms) and claims (as evidence) to a language model with a strict contract: produce an answer, cite which principle grounds each part, note caveats.
5. Validates that the answer doesn't contradict the retrieved principles by running a quick coherence check.
6. Generates an adversarial counter-position to the leaned-on principles — the strongest case against them — and includes its strongest objection as a caveat.
7. Runs a final consistency check before returning the answer.

The result is not a one-shot LLM completion; it is an answer that has been compared to the firm's existing positions and challenged by its own machinery before you see it.

3 What counts as a citation

A citation is a structured reference back to one of:

- **Claim** — an extracted statement from an upload. The most granular citation. Click-through lands on the source span.
- **Conclusion** — a synthesized firm position, possibly published as an article or as a reviewed conclusion. Citing one means the Oracle is leaning on a vetted firm output, not a raw extracted claim.
- **Principle** — a distilled, reusable position. The Oracle treats principles as soft constraints, not hard rules.
- **Upload** — a direct quote or paraphrase from an uploaded document, attached with offset and page number.

Every citation is clickable. If a citation does not click through, the Oracle has produced a broken reference — treat it as a hallucination warning even if no explicit badge is shown, and report it through **Critiques**.

4 The verbatim-citation rule

The single load-bearing rule of the Oracle (and of every other public-facing surface the workshop produces) is the *verbatim citation contract*: any quoted span the answer leans on must appear *word-for-word* in the cited source. The workshop's validator checks this on every answer. When a quoted span doesn't match, the answer is not retried with a fix-it prompt — it is rejected, and the Oracle abstains.

This is the firm's standing answer to fabrication. You will see it echoed throughout the stack: Currents abstains rather than write an opinion whose quote doesn't appear in any source; Forecasts abstains rather than publish a prediction whose reasoning cites a span that isn't there.

5 Confidence bands

The Oracle assigns each answer a band based on the strength and breadth of the citation chain:

- **high** — the answer is supported by multiple independent citations, including at least one conclusion or principle.
- **medium** — the answer is supported by at least one citation that directly addresses the question.
- **low** — citations exist but only weakly support the answer. Treat as a starting point, not as a finished position.
- **abstain** — the corpus is too thin or off-topic to answer. The Oracle returns no answer text and explains which material it would have needed.

Abstention is the correct outcome for a sparse corpus. It is not a failure. The Oracle was deliberately designed to refuse where it would otherwise fabricate.

6 The hallucination warning

A hallucination badge appears when post-hoc verification finds the answer text contains a span that the cited sources do not unambiguously support. The badge does not necessarily mean the answer is wrong; it means the support is weaker than the prose makes it sound. Standard response:

1. Read the cited sources. Decide whether the unsupported span is a paraphrase that survives or an actual overreach.
2. If it is an overreach, rephrase the question to be narrower and re-ask. The Oracle's narrower answer is usually correct.
3. If you find the Oracle consistently overreaching on a topic, log a critique — the critiques pipeline feeds back into prompt and policy tuning.

7 What to do with low-confidence answers

Low confidence does not mean “probably wrong.” It means “the corpus will not let me commit to a sharper answer.” Useful responses:

- **Add corpus.** Upload one or two documents that directly address the question. Re-ask. The same question with denser material usually moves up a band.
- **Narrow the question.** A low-confidence answer to a broad question often becomes a high-confidence answer to a more specific sub-question.
- **Promote a principle.** If the topic recurs and the firm has a working position, draft and accept a principle (Guide 2). The Oracle will start citing it.
- **Accept abstention.** For genuinely outside-corpus questions, the right answer is “the firm does not have a position on this” — the Oracle will say so.

8 The Research Advisor

Research (/research) is the other side of the Oracle: instead of answering a question, it tells you what the firm should be reading next.

It's driven by a small set of inputs:

- The firm's recent abstentions. If the Oracle keeps refusing on a topic, that topic is a candidate for a reading recommendation.
- The firm's open contradictions. Two claims pulling against each other often have a resolving source somewhere.

- Drift events — principles whose meaning has been shifting.
- The recent forecasts that have resolved against the firm (Brier score worse than baseline). A bad forecast often signals a methodology gap.

Each recommendation is a small block: a suggested title, a one- paragraph summary of why the firm should read it, a list of reading URIs (papers, articles, podcasts), and which upload or which open question prompted the suggestion.

9 The Reading Queue

Reading queue (/reading-queue) is your personal stack of suggestions. You can:

- Promote a research suggestion onto your queue.
- Mark items `reading`, `read`, or `deferred`.
- Attach quick notes that the workshop reads if you later upload the item.

The Reading Queue is intentionally low-ceremony. It's a list of things to read with light state, not a project tracker.

10 Inverse inference — “what does this event mean for our corpus?”

A specialty mode of the Oracle, available from the question form, answers the inverse question: given that some event happened, what does it *imply* about (or refute in) the firm's existing positions? It returns:

- A list of **supporting implications** — existing principles or conclusions that this event strengthens.
- A list of **refuted implications** — principles or conclusions this event undermines, with severity (mild, moderate, or severe).
- A list of **irrelevant** principles, for transparency.
- A **blindspot report** — entities, mechanisms, or adjacent topics the firm's corpus is missing in light of the event.

Severe refutations show up on the unified attention queue (Guide 6). Blindspots show up on **Research** as new reading suggestions.

11 What the Oracle will not do

- It will not answer from general world knowledge alone. If the corpus does not contain relevant material, it abstains.
- It will not generate predictions, market opinions, or trade suggestions. Those routes are gated behind the operator confirmation flow described in Guide 5.
- It will not assert a firm position that contradicts an accepted principle without an explicit flag in the answer signalling the conflict.
- It will not surface private uploads in answers visible to readers outside the firm. The public Oracle at /ask is gated to the public corpus; /codex-ask (signed in) sees the full corpus available to you.

12 Where the docs are

- `docs/methods/MQS_Specification.md` — citation contract.
- The other five guides in this series.

Where to read next

Continue with [Guide 4 — Currents](#) for the firm's outward-facing live commentary surface, which uses a related but distinct citation contract.