

What Changed in Theseus

A Plain-Language Explainer

2026-05-16

TL;DR

Theseus was rebuilt around a new idea: that what the system stores should be *principles* (generalizable truth claims) rather than *summaries* (what the author of a document seems to think). On top of that change, a few related pieces were added — a way to turn principles into structured reasoning, a cleaner contradiction engine, a way to mark each source by where it came from, an output format for the firm’s takes, and a few read views. This explainer says what each change is and why it was made — nothing about the underlying infrastructure. The final section explains how to re-process the existing corpus, which is the one thing that requires your action.

Contents

1	The core problem we fixed	2
2	What changed	2
2.1	Conclusions are now principles	2
2.2	A new layer on top of principles: algorithms	2
2.3	One contradiction engine, not six	2
2.4	Provenance is set at upload time	3
2.5	Memos as the canonical output	3
2.6	“Bet” was generalized	3
2.7	A knowledge graph view	3
2.8	Live conversation recording	3
2.9	Surface cleanup	4
3	Re-processing the existing corpus	4
3.1	Step 1: sanity-check the new extractor	4
3.2	Step 2: dry-run a small sample	4
3.3	Step 3: queue the whole corpus	4
3.4	Step 4: triage the queue	5
3.5	What it’ll cost, roughly	5
4	Where to read next, if you want depth	5

1 The core problem we fixed

The old system read a document and stored what it thought the author was saying *about themselves* — sentences like “I have since become a fan of Peter Thiel’s idea.” That kind of conclusion is autobiographical. The firm can’t act on it. We can only act on what an idea *is* and *when it applies*.

So the entire pipeline was reshaped around principles: third-person, generalizable claims about how the world works. “Definite optimism outperforms indefinite optimism over a ≥ 10 -year horizon, conditional on a coherent vision” is a principle. The firm *can* act on that — it can identify situations that match the condition and form positions about what’s likely to follow.

Everything else in this guide either follows from that change or cleans up something that was getting in the way of it.

2 What changed

2.1 Conclusions are now principles

The extractor no longer emits first-person summaries. If a source span is purely autobiographical and no underlying principle can be extracted from it, the extractor logs that fact and emits nothing. If a principle is present, it’s stored with a few extra pieces of structure: what *kind* of principle it is (rule, criterion, mechanism, heuristic, definition, formula, algorithm), what domain it applies to, what measurable quantities could test it, and a few concrete decisions it would inform.

This is the load-bearing change. The existing corpus needs to be re-processed through the new extractor — see Section 3 below.

2.2 A new layer on top of principles: algorithms

A principle says *what is true*. An algorithm says *when, given what we can observe, the principle predicts something specific*. So an algorithm is the bridge between abstract principle and concrete prediction.

Each algorithm names the observable things it watches, a condition that has to be true for it to fire, the principles it’s reasoning from, and the structured output it produces. The reasoning chain inside an algorithm cites the principles it depends on, so you can always trace a prediction back to its sources.

A canonical example: an “arms-race escalation” algorithm watches bilateral military-spending changes, rhetoric levels, and whether a third-party mediator is present. When the conditions match, it applies two principles about security dilemmas and emits a predicted spending increase with confidence bands.

Algorithms are drafted from clusters of related principles. You review each draft, accept/edit/reject it, and from then on the algorithm runs whenever its conditions are met.

2.3 One contradiction engine, not six

The old system had six different ways of detecting that two claims contradict each other — different heuristics that disagreed with each other and were hard to calibrate. That was replaced with a single detector that returns a calibrated score, a confidence band, and a human-readable explanation of where the disagreement lies.

A second piece of this change: contradictions no longer have a “resolve” button. Operator clicks were treated as authoritative, which was wrong — the database should reflect what the sources jointly imply, not which side an operator preferred on a given afternoon. Contradictions now resolve when new source material weights one side decisively over the other.

2.4 Provenance is set at upload time

Every piece of source material now carries one of four labels, chosen when it's uploaded:

- **Proprietary** — the firm wrote it.
- **Endorsed external** — someone else wrote it, but we explicitly endorse it as representative of our thinking.
- **Studied external** — reference material; we read it but don't endorse it.
- **Opposing external** — material we explicitly disagree with, kept for the value of testing our positions against it.

Why this matters: the system used to flag contradictions between the firm's principles and opposing material it was reading — which is noise, because we *expect* to disagree with that material. Provenance demarcation lets the contradiction engine skip those cross-provenance pairs and surface only the contradictions that are actually news.

2.5 Memos as the canonical output

When the system answers a question — whether it's "what's our take on X?" or "should we bet on Y?" — the answer now takes a fixed memo shape: a TL;DR, the question, the governing principles, the observed inputs, the reasoning chain, the implied position, what would change our mind, and any caveats or abstentions. The memo is the firm's canonical output.

The system explicitly prefers to *abstain* over making a chain of reasoning it can't ground. If there are no governing principles, or the principles directly contradict each other, or the confidence band is too wide, or the question itself is unformed — the memo says so, by name.

2.6 "Bet" was generalized

A bet used to mean a financial position on a prediction market or an equity. "Bet" now means any falsifiable commitment of firm resources: a financial position on a market, a public statement of position with no money behind it, an internal allocation of operator time or hiring direction, or a scientific prediction that resolves against external data. All four are tracked in one place.

The reason: the firm's edge is the principle layer, not any specific way of expressing a view. Restricting bets to financial markets undersold what the system could do.

2.7 A knowledge graph view

The corpus is no longer a flat list of claims. There's now a cross-source graph view where principles, sources, algorithms, memos, and concepts are nodes, and the relationships between them are edges ("derived from", "contradicts", "supports", "applies to", "cites"). Click any edge and a reasoner produces a grounded explanation of why that edge exists. The reasoner refuses to fabricate connections the data doesn't support.

This is the reading view for "what does the corpus jointly imply about X?" — a question that wasn't directly answerable before, because the underlying structure didn't exist.

2.8 Live conversation recording

The meeting-recording surface ("Dialectic") now fires real-time contradiction flags as a participant says something that contradicts (a) something they or another participant said earlier in the same conversation, or (b) a principle the firm has already committed to. Provisional principles can be extracted live and triaged after the session ends.

The motivation: the most valuable corpus material is often the live conversation, but it was historically lost or transcribed only as an afterthought. Real-time flags turn the conversation itself into a feedback loop with the firm's existing commitments.

2.9 Surface cleanup

A few smaller changes worth knowing about: the homepage and about page now lead with the "Philosopher in a Box" positioning — Theseus is the firm's machine, not a product for sale. The old "Attention" panel on the dashboard is gone (redundant once the Currents feed and algorithm-invocation feeds matured). The legacy contradiction heuristics are marked deprecated and no longer write new rows.

3 Re-processing the existing corpus

TL;DR

This is the one thing that requires your action. Existing conclusions are NOT auto-replaced — the new extractor's output goes to a triage queue at / (authed) /extractor/re-extract where you accept, edit, or reject each replacement. Do a small sample first to confirm the new extractor behaves as you want.

3.1 Step 1: sanity-check the new extractor

Before queueing the whole corpus, run the new extractor against one or two artifacts whose existing conclusions had the first-person symptom (e.g., the one that produced "I have since become a fan of Peter Thiel's idea"). The single-artifact dry-run prints the proposed principles without persisting anything:

```
cd ~/Desktop/Theseus
source .venv-currents/bin/activate
python noosphere/scripts/reextract_corpus.py \
    --artifact <artifact_id> --dry-run
```

Confirm the output: each proposed principle should be third-person and generalizable, the principle-kind should match your reading, the source citation should be intact, and refusals should fire on spans that are purely autobiographical.

3.2 Step 2: dry-run a small sample

Once one or two artifacts look right, sample five at random to confirm the extractor is consistent:

```
python noosphere/scripts/reextract_corpus.py --sample 5 --dry-run
```

If any of the five looks wrong, stop and adjust the principle extractor's system prompt before going further.

3.3 Step 3: queue the whole corpus

When the sample looks right, queue every artifact for re-extraction. You can do this all at once, or restrict by date or by provenance:

```
# Everything:
python noosphere/scripts/reextract_corpus.py --all

# Only your own writing first (highest signal, lowest noise):
python noosphere/scripts/reextract_corpus.py --all \
```

```
--provenance PROPRIETARY

# Only material added since 2026-03-01:
python noosphere/scripts/reextract_corpus.py --all \
    --since 2026-03-01
```

The job is asynchronous and respects the system's hourly budget, so a large corpus may take several cycles to drain.

3.4 Step 4: triage the queue

Each proposed replacement lands at / (authed) /extractor/re-extract, NOT as an automatic overwrite. You review each row, which shows:

- The original source span.
- The existing first-person conclusion (from the old extractor).
- The new principle-shaped proposal.
- Three buttons: **Accept**, **Edit & Accept**, or **Reject**.

The system never replaces anything on its own — every row needs your confirmation. Once you accept a row, the new principle gets embedded, joins a cluster, gets tested for contradictions against its siblings, and becomes searchable on the principles page. Algorithms may be drafted from new clusters automatically.

3.5 What it'll cost, roughly

For a corpus of ~500 artifacts:

- **Tokens.** Roughly 1.5–2M total tokens.
- **Wall clock.** At the default hourly budget, expect 6–24 hours for the queue to drain (smaller corpora proportional).
- **Your triage time.** Each row takes 30–60 seconds if the proposed principle is reasonable. For 500 artifacts, budget 4–8 hours of triage, spread across whichever days work. There's no rush.

4 Where to read next, if you want depth

The earlier guides in this series go deeper on individual surfaces:

- docs/guides/01_Theseus_Quick_Start.pdf — the minimal intro for a new reader.
- docs/guides/02_Knowledge_and_Principles.pdf — how ingestion-to-principles works.
- docs/guides/03_The_Oracle.pdf — using the Oracle.
- docs/guides/04_Currents.pdf — the live opinion feed.
- docs/guides/05_Forecasts_and_Portfolio.pdf — prediction markets and equities.
- docs/guides/06_Operator_Console.pdf — the operator surface and the safety architecture in detail.