

Adversarial probing of hidden assumptions: a firm cluster

Cluster adversarial-audit-2a74eb9b

Theseus Auto-Paper Generator

May 2026

Disclosure: machine-drafted, founder-reviewed.

Lead conclusion: **conclusion-adversarial-audit-1**.

Methodology root: **profile-adversarial-audit-1** (adversarial-audit).

Abstract

This paper distills 2 firm conclusion(s) anchored on conclusion-adversarial-audit-1, all sharing the methodology root "Adversarial probing of hidden assumptions". The cluster's claims are stress-tested against 1 resolved forecast(s) drawn from the firm's calibration ledger. Lead claim: Adversarial review surfaces the hidden assumption a friendly read leaves buried; we will block publication if an objection is unaddressed.

1 Introduction

The Theseus engine routinely produces firm conclusions, attaches them to a methodology profile, and tracks how those conclusions perform when their predictive consequences are settled in public forecast markets. This paper presents one such cluster of 2 conclusion(s), drawn together by their shared methodology root and by at least one resolved forecast that touches the cluster. The cluster's lead claim states: "Adversarial review surfaces the hidden assumption a friendly read leaves buried; we will block publication if an objection is unaddressed."

Conclusions in this cluster

- **conclusion-adversarial-audit-1** [row:conclusion:conclusion-adversarial-audit-1] — Adversarial review surfaces the hidden assumption a friendly read leaves buried; we will block publication if an objection is unaddressed.
- **conclusion-adversarial-audit-2** [row:conclusion:conclusion-adversarial-audit-2] — A buried assumption stays buried under a friendly read, so the firm runs the adversarial reviewer first.

2 Methods

The cluster shares a single methodology root, *Adversarial probing of hidden assumptions* (profile-adversarial-audit-pattern *adversarial-audit*).

Summary. Surface the assumption a friendly read leaves buried by running an adversarial reviewer against every load-bearing claim before it is published.

Reasoning moves

- ask what would have to be true for the claim to fail
- run the adversarial reviewer before the friendly one
- treat an unaddressed objection as a publication blocker

Assumptions

- the adversarial reviewer is not subject to the same blind spot
- objections are cheaper to surface than to discover post-hoc
- a buried assumption is load-bearing until shown otherwise

3 Results

The cluster’s conclusions were stress-tested against 1 resolved forecast(s) from the firm’s calibration ledger. Each row below resolves to a database record.

Headline	Outcome	Stated	Brier	Row
Adversarial probing of hidden assumptions: claim 1 holds	YES	0.70	0.090	[row:forecast_resolution:prediction]

Across the 1 fully-scored forecast(s) above, the firm achieved a mean Brier score of 0.090. Each contributing row is auditable via its \rowref{...} marker.

4 Discussion

Reviewer firm-peer-review on conclusion-adversarial-audit-1: verdict revise (confidence 0.78).

Peer-review findings

- **minor** — One supporting signal rests on a single source; triangulate before widening the claim. [row:review_report:review-adversarial-audit-seed]

5 Limits and failure modes

The methodology root carries a registered failure-mode catalog. Every limit below is an entry from the firm’s catalog, not a synthesized observation.

- adversarial monoculture shares the friendly reviewer’s blind spot

- objection fatigue downgrades a structural objection to a nitpick

6 References

1. Firm source memo: Adversarial probing of hidden assumptions, May 2026 [row:artifact:artifact-adversarial]

This paper was machine-drafted by Theseus from the firm's existing artifacts. Every numerical claim resolves to a row in the firm's database (see [row:...] markers). The .tex source is the authoritative artifact; the PDF is a build product.