

Theseus Seasonal Research Review

2026 Q2

Theseus

2026-06-30 12:00 UTC

Disclosure: machine-drafted, founder-reviewed.

Slug: 2026_Q2_Review. Generated 2026-06-30 12:00 UTC.

Sections marked *data not available* are omitted from the quarter rather than estimated.

Overview

The firm closed 2026 Q2 with 4 active method(s), 1 deprecated, and 1 retired. 2 drift event(s) entered the ledger; 4 forecast(s) resolved. The firm published 2 article(s) and the self-critique pass surfaced 2 finding(s).

1 Methods

The register carries 4 active method(s) at quarter close. The firm marked 1 method(s) deprecated and retired 1. Retirement is the firm's standing acknowledgement that a method no longer earns its place in the register. **Active.** 4 method(s). **Deprecated.** 1 method(s). **Retired.** 1 method(s).

Retired this register

- `m_naive_priors` — Naive priors (*v1.0*)

Deprecated

- `m_legacy_heuristic` — Legacy domain heuristic (*v1.0*)

2 Drift

The firm observed 2 drift event(s) in window. The highest-scoring event was against target `m_geometry` at score 0.420; the firm treats the score as evidence of a regime shift to be investigated, not an automatic retirement trigger. **Drift events in window: 2.**

Target	Score	Observed	Notes
<code>m_geometry</code>	0.420	2026-04-21	cross-domain transfer regression after embedder swap
<code>m_calibration</code>	0.275	2026-05-04	bucket-12 reliability slipped

3 Calibration trend

Across 4 resolved forecast(s), the firm posted a mean Brier of 0.138, and mean log-loss 0.340. The firm reads the score as a calibration check, not a performance leaderboard. Resolved forecasts in window: 4. Mean Brier: 0.138. Mean log-loss: 0.340.

4 Open questions

data not available

5 Published articles

The firm published 2 article(s) in the quarter. Each carries the machine-drafted, founder-reviewed disclosure non-removably; the firm speaks as the firm, never as a single voice.

- 2026-04-15 — A calibrated narrow claim beats a confident broad one (**calibrated-narrowing**)
- 2026-05-03 — Adversarial review surfaces the buried assumption (**adversarial-first**)

6 Principles distilled

The firm distilled 2 principle(s) in the quarter; the strongest reached conviction 0.74 across 3 domain(s). Conviction earns its place in the register; the firm does not auto-promote a draft.

- Cross-domain coherence beats single-domain depth when the firm must commit under uncertainty. (*domains: 3, conviction: 0.74*)
- Adversarial review precedes friendly review; the friendly read leaves load-bearing assumptions buried. (*domains: 2, conviction: 0.68*)

7 Most-edited conclusions

no conclusions edited in window

8 What we got wrong

The self-critique pass surfaced 2 finding(s) the firm got wrong this quarter. Each finding is named — the firm does not soften a self-critique.

- Article **calibrated-narrowing** — Self-critique on 'A calibrated narrow claim beats a confident broad one' — verdict: weakened. The supporting forecast resolved YES, but the claim's narrowing was tighter than the resolved bucket licensed. (*review item review_self_crit_one*)
- Article **adversarial-first** — Self-critique on 'Adversarial review surfaces the buried assumption' — verdict: scope-overreach. The supporting evidence covered methodology and forecasting, not all transfer targets the article implied. (*review item review_self_crit_two*)

*This review was machine-assembled from the firm's own machinery for 2026 Q2. The .tex source is the authoritative artifact; the PDF is a build product. The structured numbers carried in the sibling **review.json** are the source of truth for every numeric claim above.*