

The Theseus Stack

A plain-language companion guide

(Dialectic · Noosphere · Theseus Codex)

Read this guide first.

Then read the technical guide.

What this is. A friendly introduction to the three programs Theseus has built. It uses the same section order as the more detailed technical guide, so once you have the big picture from here you can “upgrade” to that document section by section. No formulas, no module names, no code paths — just the ideas and the way they fit together.

Who it’s for. Anyone who wants to understand what Theseus actually *is* without having to first learn the vocabulary of the firm. If you are familiar with the firm’s research — coherence, contradiction geometry, methodology — you may want to skip straight to the technical guide.

Contents

1	What Theseus is, and why these three programs exist	2
1.1	The firm in one sentence	2
1.2	What each piece does, in one paragraph	2
1.3	A single picture	3
2	Dialectic: the smart notetaker	3
2.1	What it feels like to use it	3
2.2	What it’s actually doing, behind the scenes	4
2.3	The polite assistant (“Theseus speaks”)	4
2.4	What gets saved at the end of a session	5
3	Noosphere: the workshop	5
3.1	What Noosphere is for	5
3.2	What happens to a piece of material, step by step	5
3.3	The six coherence checks (without the math)	6
3.4	Adversarial objections (the “Dialectical Crucible”)	7
3.5	The live pipelines: Currents and Forecasts	7
3.6	The hard rules	8
4	Theseus Codex: the dashboard and the website	8
4.1	Two faces of one website	8
4.2	What happens when you upload a file	9
4.3	From conclusion to public article	10
4.4	Some careful design choices	10
4.5	The desktop version	10
5	How the three pieces fit together	10
5.1	One story, from microphone to public article	11
5.2	What the three pieces have in common	11
6	The bigger idea	12
6.1	Why the firm built this rather than just writing essays	12
6.2	Five questions the system asks about every conclusion	12
6.3	Six things every published conclusion must carry	12
6.4	The places the system is designed to abstain	13
6.5	Where to go next	13

1 What Theseus is, and why these three programs exist

1.1 The firm in one sentence

Theseus is a small research and investment firm whose actual product isn't its opinions — it's the *way* those opinions get formed. Anyone can have an opinion about a company, a market, or an idea. What's rare is being able to show, with receipts, exactly how the opinion was arrived at, what was considered, what was rejected, who disagreed, and what would have to change for the opinion to change.

The three programs are tools for keeping that kind of receipt.

1.2 What each piece does, in one paragraph

Dialectic — the smart notetaker

Imagine a meeting where, an hour in, somebody says “we should ship as fast as possible,” and then half an hour later somebody else says “we should never ship without three reviews.” No one notices the clash. Dialectic is a desktop program that listens to the meeting and catches things like that. It transcribes what people say, picks out the actual claims being made, and flashes a red card on screen when two claims pull against each other. After the session it saves a tidy file the rest of the system can read.


Noosphere — the workshop

Now imagine all the firm's raw material in one room: podcast recordings, essays, PDFs, tweets, prediction-market quotes. Noosphere is the workshop where that material is taken apart. It pulls out each individual claim, groups similar claims into more durable “principles,” and builds a map of how those principles relate (one supports another, one contradicts another, one refines another). It runs six different checks for whether ideas hold together. It automatically generates the strongest objection it can find against each new conclusion. And it does this on an ongoing basis, including for current events and prediction markets.

Theseus Codex — the dashboard and the website

Finally, all of that work needs somewhere to live and somewhere to be shown. The Codex is a website. Founders log in to upload material, review what the workshop has produced, and decide what's ready for the world. The same site, viewed without logging in, is the firm's public face: published conclusions, articles, opinions on current events, the prediction-market track record, a critique submission form, and so on.

1.3 A single picture



The way to read this diagram: humans and the world create raw material at the top; the middle row does the work; and the website at the bottom is where everything lives and where it gets shown.

2 Dialectic: the smart notetaker

2.1 What it feels like to use it

You open Dialectic before a working session, sign in to your firm account, and click `BEGIN SESSION`. From there, you just talk. You don't have to do anything else — the program is listening through the microphone.

On screen you see four things at once:

- A running transcript of what's being said.
- A little graph of “claims” as labeled dots. Every meaningful sentence becomes a dot.
- A panel called `TENSIONS` where, if two of those dots clash, the program puts up a red-bordered card: “this earlier thing and this later thing don't fit together.”
- A panel called `PROMPTS`, where (if you've turned it on) the program can politely interject with a follow-up question.

When you click `STOP`, the program saves a tidy file of everything it analyzed. The original audio stays on your laptop. The analysis file is the one that gets sent up to the firm's website.

2.2 What it’s actually doing, behind the scenes

In plain words

Think of the program as a chain of small workers, each doing one job and handing off to the next.

1. One worker listens to the microphone and chops the sound into pieces of speech, ignoring silence.
2. A second worker turns each piece of speech into text.
3. A third worker reads the text and decides what’s a *claim* (an actual assertion someone is making) versus what’s filler, hedging, or a question.
4. A fourth worker turns each claim into a kind of mathematical fingerprint, so that two claims with similar meaning end up with similar fingerprints.
5. A fifth worker compares every new claim’s fingerprint to the recent ones to see what topic it belongs to.
6. A sixth worker runs each new claim against the recent claims to see if any of them clash. When two clash hard enough, that’s a contradiction.
7. A final worker writes the result to disk and updates the screen.

This is all happening in parallel, and the screen only redraws once a second. The point is that none of the heavy work happens on the part of the program that draws the window, so it stays responsive while everything else hums along.

2.3 The polite assistant (“Theseus speaks”)

You can optionally turn on an **interlocutor** mode, which is the program speaking up during the session. By default it stays silent. If you turn it on, it can do three kinds of things:

- Flag a contradiction it just found.
- Point out a question that was raised earlier and hasn’t been addressed.
- Note a prediction someone just made, so you can revisit it later.

It speaks in third person (“Theseus notes that...”) — on purpose, so it doesn’t pretend to be a person in the room. It has a strict budget for how often it can speak per session, a cooldown between interventions, and you can say “Theseus, stand down” out loud at any time to make it stop.

Every intervention it considers — whether it actually spoke up or stayed quiet — is logged. After the session, you can review all of them in the firm’s dashboard.

2.4 What gets saved at the end of a session

- The audio file, kept on your laptop only.
- The analysis file — one line per claim, with the time, who said it, what was said, its fingerprint, which other claims it clashed with, and what topic it belongs to.
- An intervention log, if the assistant was turned on.
- A short “reflection” summary of the session.

If your firm credentials are set up, the analysis file and the reflection are automatically uploaded to the firm’s website right after you stop. The audio is never uploaded automatically; that stays with you.

3 Noosphere: the workshop

3.1 What Noosphere is for

Anything the firm cares about eventually goes through Noosphere. A 4-hour podcast becomes claims and principles. An essay becomes claims and principles. A PDF becomes claims and principles. A tweet might become a current-event opinion. A prediction-market listing might become a probability forecast. The shape of the workshop is roughly the same every time.

3.2 What happens to a piece of material, step by step

1. Read the source.

- If it’s an audio file, transcribe it.
- If it’s a PDF, pull the text out (and if the PDF is just scanned images, run the OCR step).
- If it’s text, just read it.

2. Strip out things that aren’t the speaker’s own.

For example, if a podcast has an interview prompt at the start (“what do you think of X?”), that’s the host asking, not the founder asserting. The program removes it before analysis. Similarly for quoted positions from other thinkers.

3. Break the material into claims.

A claim is one sentence, one speaker, one assertion. The program doesn’t summarize — it preserves the individual sentences, because the whole point is to be able to reason about them one at a time.

4. Group similar claims.

If the founder makes the same point in four different ways over four months, those should all roll up to one durable “principle.” The program does this by comparing the fingerprints of claims and clustering similar ones together.

5. **Build a map.** Principles get put into a network, and the connections between them are labeled — this one *supports* that one, this one *contradicts* that one, this one *refines* that one, and so on.
6. **Run coherence checks.** For any two claims (or principles), the program asks six different versions of the question “do these two go together?” — and only treats the answer as confident when four of the six agree (see below).
7. **Profile the reasoning.** The program also notes *how* a piece of material reasoned — did it work from first principles? Did it transfer ideas across domains? Did it adversarially revise its own position? Each pattern is one row in a separate table.
8. **Generate objections.** For each meaningful new conclusion, the program automatically tries to construct the strongest possible objections to it, from three different intellectual traditions. If an objection is convincing enough, the conclusion’s confidence gets demoted.
9. **Track over time.** Every time a principle is reinforced, it records when, by whom, and whether its meaning has drifted. A principle that hasn’t been mentioned in a while loses conviction; a principle whose meaning has subtly changed is flagged for review.
10. **Hand everything back to the website.** The Codex is the place humans look at the results, so the workshop writes its output back into the website’s database.

3.3 The six coherence checks (without the math)

This is the heart of the workshop. When the system wants to know whether two claims fit together — or actively conflict — it runs six different tests:

1. **A language test.** Does this claim entail that one? Contradict it? Or are they unrelated?
2. **An argument-shape test.** If you drew the two claims and the relationship between them as a little argument diagram, would both sides survive?
3. **A probability test.** Do the probabilities the claims imply add up sensibly, or do they break basic arithmetic?
4. **A geometry test.** If you look at the meaning “fingerprints” of the two claims, does their difference look like the kind of pattern that contradiction usually produces?
5. **A compression test.** If you stick the two claims together, can the result be compressed more tightly than the two on their own? (Shared structure compresses; disagreement doesn’t.)
6. **A reasoned check.** A final pass where a model is asked for its verdict — but is forced to name and quote the first five tests’ results in its explanation, so it can’t quietly hand-wave.

Each test votes. If four or more agree, that’s the verdict (*coherent*, *contradicting*, or *unresolved*). Otherwise the pair gets parked in “unresolved” for later human review.

Why six rather than one

A single test would be too easy to fool — different kinds of disagreement leave different fingerprints. “The sky is blue” and “the sky is not blue” look almost identical to most language models if you just compare them. But the way their meanings *differ* from each other is unusually narrow, and that narrowness is what the geometry test picks up. The six tests together cover blind spots no single test has.

3.4 Adversarial objections (the “Dialectical Crucible”)

This is the part that gives the firm its character. When a new conclusion lands, the program doesn’t just file it away. It actively tries to attack it.

1. It builds a record of what the firm has already said on the topic, so it can tell genuine novelty from re-litigation.
2. It asks a language model to generate three objections, each from a distinct intellectual tradition, with specific atomic claims attached.
3. Each objection gets run through the same six coherence checks against the conclusion.
4. If any objection wins those checks with high enough confidence, the conclusion is automatically demoted to a lower confidence tier.

In other words: a conclusion that can’t withstand the strongest objection the system can generate doesn’t keep its rating. This is enforcement, not commentary.

3.5 The live pipelines: Currents and Forecasts

The same workshop also runs two ongoing pipelines.

Currents watches social media and news for events the firm might want to comment on. For each event, it asks: do we have at least two solid, on-topic conclusions to ground an opinion in? If yes, it drafts an opinion (with each quoted source span verified to literally appear in the source), saves it, and over time clusters opinions into public articles. If not, it abstains, with an explicit reason like “insufficient sources” or “off domain.”

Forecasts watches Polymarket and Kalshi for prediction markets relevant to the firm’s beliefs. For each market, it retrieves the relevant material, runs a deterministic decision procedure, and (if it has enough material and is far enough from the market’s close) generates a probability prediction with citations. If certain conditions are met, it also stages a paper bet. Live bets require explicit per-prediction authorization from a founder — this is enforced at the database level. When markets resolve, the system scores itself (Brier score, log-loss) and updates a public calibration page.

3.6 The hard rules

Several rules are non-negotiable throughout the workshop:

- **Quote-or-abstain.** Every quoted source span must appear word-for-word in the source. If it doesn't, the system abstains rather than retrying.
- **Cite-the-receipts judge.** The reasoned coherence check can't override the other five unless it cites them by name and value.
- **Wrap every prompt.** Every text fed to a language model is wrapped in unambiguous markers so the model can't be misled by hostile material in the source.
- **Version everything.** When the firm updates how it processes material, old conclusions don't get wiped — they get flagged for re-review.

4 Theseus Codex: the dashboard and the website

4.1 Two faces of one website

The Codex is a normal website. When you visit it without logging in, you see the firm's public face: reviewed conclusions, articles, opinions on current events, prediction-market track record, methodology pages, an "ask the firm" question form, a critique submission form. Subscribers can sign up by email and get a periodic digest — no tracking pixels, just a clickable acknowledgment if they want to confirm they read it.

When a founder logs in, the same website transforms into a private dashboard. There are around seventy founder-only pages. The main ones:

- **Dashboard.** The home page. Shows the firm's active principles (ranked by conviction), recent conclusions, recent uploads, drift events, and operational health (is the current-events pipeline running? is the embedding index healthy? where does the portfolio stand?).
- **Library.** An inventory of every upload across the firm.
- **Upload.** The form for adding new material, plus a per-upload page that streams the workshop's progress live as it processes the new file.
- **Transcripts.** A reader view of any uploaded transcript, with speakers and timestamps.
- **Conclusions.** A browseable list of every conclusion, with their confidence tier, sources, methodology profile, lineage links, and adversarial status.
- **Contradictions and open questions.** Things the system flagged as unresolved, for the founder to address.

- **Explorer.** A 3D map of the firm’s conclusions, where each conclusion is a point in space and similar conclusions cluster together. You can fly around it.
- **Founder Currents and Forecasts.** The founder-side views of the live pipelines, including operator controls (kill switch, live-bet authorization).
- **Publication queue.** The gate before anything goes public. A founder reviews a draft, checks a small list (was the methodology okay? has it been challenged? is the writing clear? did anything private leak?), and either approves, sends back for revision, or declines.
- **Principles triage.** Newly distilled principles wait here until a founder accepts, rejects, or merges them.
- **Methods.** A browser for the firm’s registry of reasoning methods, with a per-method page showing track record, failure modes, and domain of applicability.
- **Methodology review week.** A quarterly five-day process where founders walk through the firm’s recent work, with each day cryptographically signed once written.
- **Account and API keys.** Where founders mint the API keys that the Dialectic desktop app uses to log in.
- **Operations.** Dashboards for system health, retention, and CI status.

4.2 What happens when you upload a file

1. The website checks your credentials (either a logged-in session or an API key from the Dialectic app).
2. It verifies the file isn’t too big and is a kind of file it knows about.
3. It extracts the text from the file — straight read for text, a PDF library for PDFs, a Word library for `.docx`, the cloud speech-to-text service for audio.
4. It saves the new upload to the database with a status of “pending.”
5. It sends back an immediate acknowledgment to the user, then in the background tries one or more of three paths to actually process the file: spawn the workshop locally, drop the job into a queue for a worker to pick up, or kick off a workflow on GitHub’s servers — whichever fits the deployment.
6. As the workshop runs, its output is streamed live into the upload’s row in the database, so the founder can refresh the page and watch the progress.
7. When the workshop is done, the upload’s status flips to “ingested” and the new conclusions, contradictions, and methodology profiles appear on the dashboard.

4.3 From conclusion to public article

1. A new conclusion appears on the dashboard.
2. A founder opens its page, reviews the sources, the methodology profile, the contradictions raised, and the adversarial objections.
3. If publishable, the founder sends it to the publication queue.
4. In the queue, it gets a four-item checklist (methodology okay, adversarial engagement okay, clarity okay, no leakage of private material).
5. On approval, the website mints a permanent public version (with its own URL and version number), and the workshop signs it cryptographically so the public version can later be checked against tampering.
6. The public page goes live. Subscribers get notified via the digest.

Citations on public pages are gated to public sources only — private material is never exposed through the public side, even when a conclusion was originally grounded in something private.

4.4 Some careful design choices

- **There’s no signup form.** Founders are added by admins. The firm is small and the bar for being inside the system is deliberately high.
- **Founder accounts are never hard-deleted.** If a founder leaves, their record stays so the audit trail of past conclusions, replies, and confirmations stays intact.
- **The signing keys never live on the website.** Cryptographic signatures are written by the workshop and only stored by the website.
- **Reader emails never have tracking pixels.** If a subscriber wants to confirm they read a digest, they click a link; otherwise the firm has no idea.
- **The website never sends money.** If an outside critique earns a bounty, the system queues an intent that a founder has to confirm by hand.

4.5 The desktop version

There’s also a desktop app version of the Codex — the same website, packaged with its own local database, so a founder can use it offline. It auto-updates from the same source as the live site. Internally it just runs the website locally and shows it in a window.

5 How the three pieces fit together

5.1 One story, from microphone to public article

1. Two founders sit down for a 4-hour podcast. **Dialectic** is running on a laptop. It transcribes the discussion, picks out claims, catches three contradictions in the moment, prompts one follow-up question, and saves the session when they stop.
2. Dialectic uploads the analyzed session to the Codex website.
3. The Codex receives the upload, creates a new entry, and hands the file to **Noosphere**.
4. Noosphere chunks the transcript, extracts atomic claims, embeds them, distills new principles, builds the knowledge graph, runs the six coherence checks across the new material, generates objections to the new conclusions, profiles the reasoning patterns, and writes the results back into the Codex’s database.
5. The founders open the Codex dashboard the next morning. They see new conclusions, contradictions to resolve, and a research advisor’s suggestion of what to read next.
6. One conclusion is strong, well-sourced, and has survived the adversarial pass. They push it to the publication queue.
7. They review it against the four-item checklist and approve it. The Codex mints a permanent version, and Noosphere signs it.
8. The conclusion goes public. Subscribers get a digest mention.
9. The next day, a relevant tweet appears in the world. Noosphere’s Currents pipeline picks it up, finds at least two relevant firm conclusions to ground an opinion in, drafts the opinion, verifies every quoted span literally appears in the cited sources, and saves it. A few hours later, with several related opinions in hand, an article-draft is generated and queued.
10. A relevant Polymarket market opens. The Forecasts pipeline retrieves the same material, builds a deterministic decision trace, generates a probability prediction with citations, and (because the conditions for a paper bet are met) stages one.
11. Weeks later, the firm updates the way it extracts methodology profiles. A founder runs the reanalysis script. Every existing upload is re-processed in place, and old conclusions are updated rather than wiped.

5.2 What the three pieces have in common

- **A shared notion of “claim.”** The same simple shape — one sentence, one speaker, one fingerprint — shows up in Dialectic, in Noosphere, and in the Codex’s database.
- **The quote-or-abstain rule.** Every quoted source span has to appear in the source. The system always prefers to say “I don’t know” than to make something up.

- **Versioned methods.** When the firm changes how it reasons, the change is recorded with a version number, and old results can be re-evaluated against the new version on demand.
- **Cryptographic signatures owned by the workshop.** Public artifacts are signed; the website only stores the signatures.

6 The bigger idea

6.1 Why the firm built this rather than just writing essays

Most research firms produce conclusions. Their value is their opinion. The problem is that an opinion, by itself, is hard to evaluate from the outside. Did they get to it for good reasons? Did they consider what would prove them wrong? Did they update when the facts changed?

Theseus’s bet is that, at scale, the more durable thing is not the opinion but the *way of arriving at the opinion*. So the firm builds software whose job is to keep the record honest: what was read, what was concluded, by what method, against what objections, under what conditions for revision.

If the bet is right, then over time, a firm that publishes opinions *with their receipts* should be more trustworthy than one that just publishes opinions.

6.2 Five questions the system asks about every conclusion

1. **Did the method actually generate a claim that could have been wrong?** (A method that always agrees with itself isn’t useful.)
2. **Was the test severe enough that the claim would have failed if it were wrong?** (A test that passes everything is no test.)
3. **Does the method actually fit the question?** (A great valuation screen does not, by itself, make a great hiring decision.)
4. **How many special exceptions does the conclusion need to work?** (A claim that needs lots of caveats is fragile.)
5. **Where should this method stop being trusted?** (No reasoning method works everywhere.)

The Codex puts a number on each of these for every conclusion, and combines them into a quality score — with the rule that a weak “where does this stop working?” answer can’t be papered over by strong scores on the other four.

6.3 Six things every published conclusion must carry

1. The conclusion itself.

2. The method that produced it.
3. The trail of sources behind it.
4. The strongest objection encountered (and whether it survived).
5. The condition under which it would be revised.
6. A way to come back and check, after some time, whether the world later confirmed or refuted it.

Each of these maps to something concrete in the system. If any of the six is missing, the conclusion structurally cannot be published.

6.4 The places the system is designed to abstain

Several places in the software are deliberately set up to say “no” rather than push through:

- If a current event doesn’t have at least two solid, on-topic firm sources, the system publicly abstains rather than improvising an opinion.
- If a quoted span doesn’t appear word-for-word in a source, the system refuses to use it.
- If a prediction market is about to close, the system refuses to predict on it.
- If a public-article check finds that private material would leak, the publication is blocked.
- If an objection lands strongly enough against an existing conclusion, the conclusion gets demoted automatically — the system doesn’t wait for someone to argue about it.
- Live bets cannot be placed without explicit per-prediction authorization, enforced inside the database itself.

6.5 Where to go next

Once you’ve internalized the shape of all this, the technical guide (the document this guide is a companion to) walks the same six sections again, but with the algorithms, formulas, file paths, and module names that make the system actually run. Read it section by section — you should find that each section in the technical guide has a counterpart here, so the “upgrade” is roughly: take any section of this guide, and read the corresponding section of the technical guide for the wiring underneath.

This guide deliberately omits formulas, file paths, and most precise numbers. If you find yourself wanting them, that’s the right time to move to the technical guide.